



Identifying a Whale by it's Tail

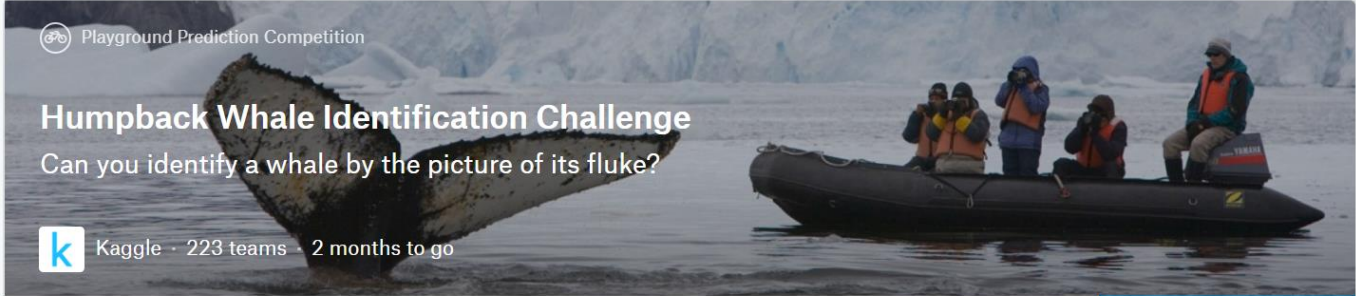
COMPARISON OF VARIOUS IMAGE CLASSIFICATION METHODS

About the Dataset

FROM KAGGLE COMPETITION

9000+ IMAGES

1900+ GROUPS



Playground Prediction Competition


Humpback Whale Identification Challenge

Can you identify a whale by the picture of its fluke?

Kaggle · 223 teams · 2 months to go

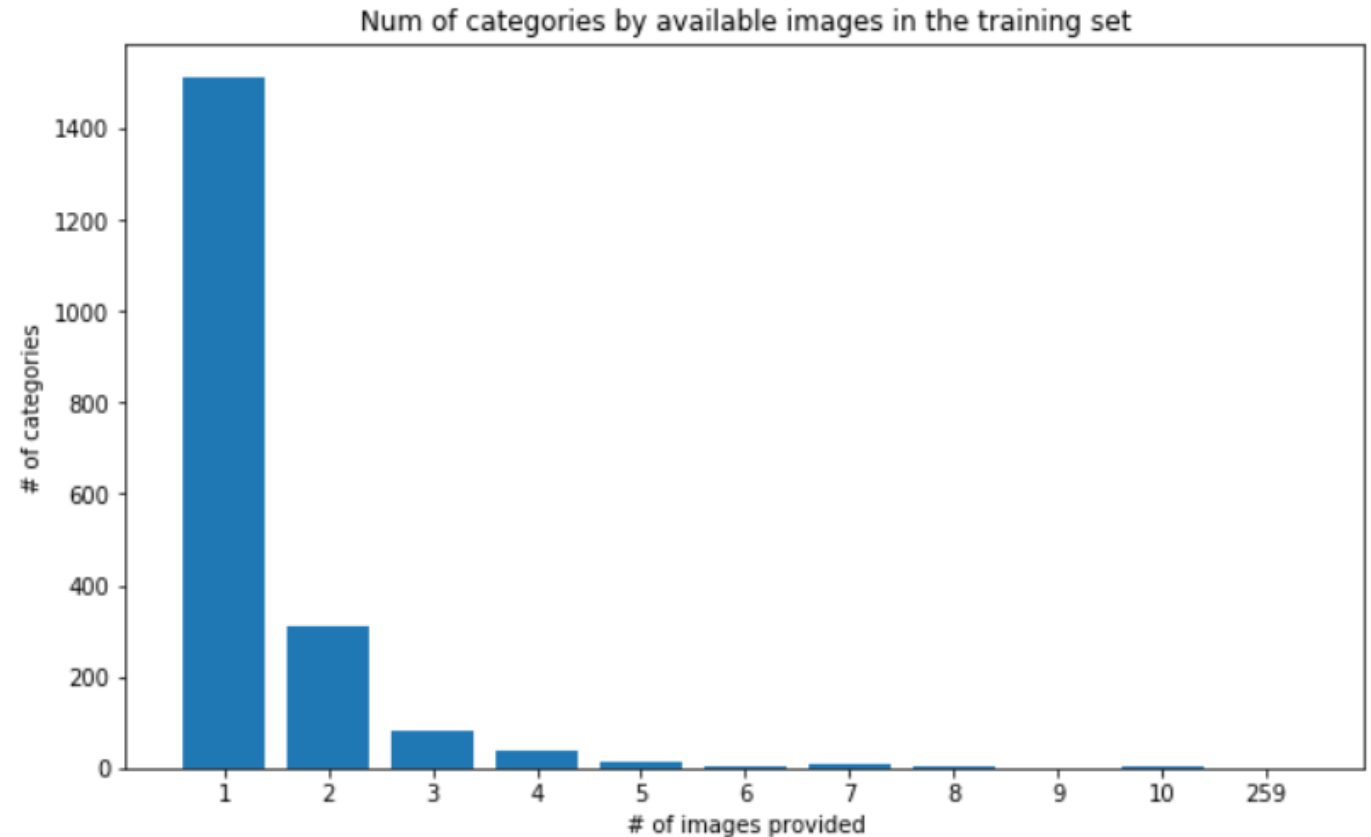
[Overview](#) [Data](#) [Kernels](#) [Discussion](#) [Leaderboard](#) [Rules](#) [Team](#) [My Submissions](#) [Submit Predictions](#)

Overview

Description	<p>After centuries of intense whaling, recovering whale populations still have a hard time adapting to warming oceans and struggle to compete every day with the industrial fishing industry for food.</p> <p>To aid whale conservation efforts, scientists use photo surveillance systems to monitor ocean activity. They use the</p>	
Evaluation		

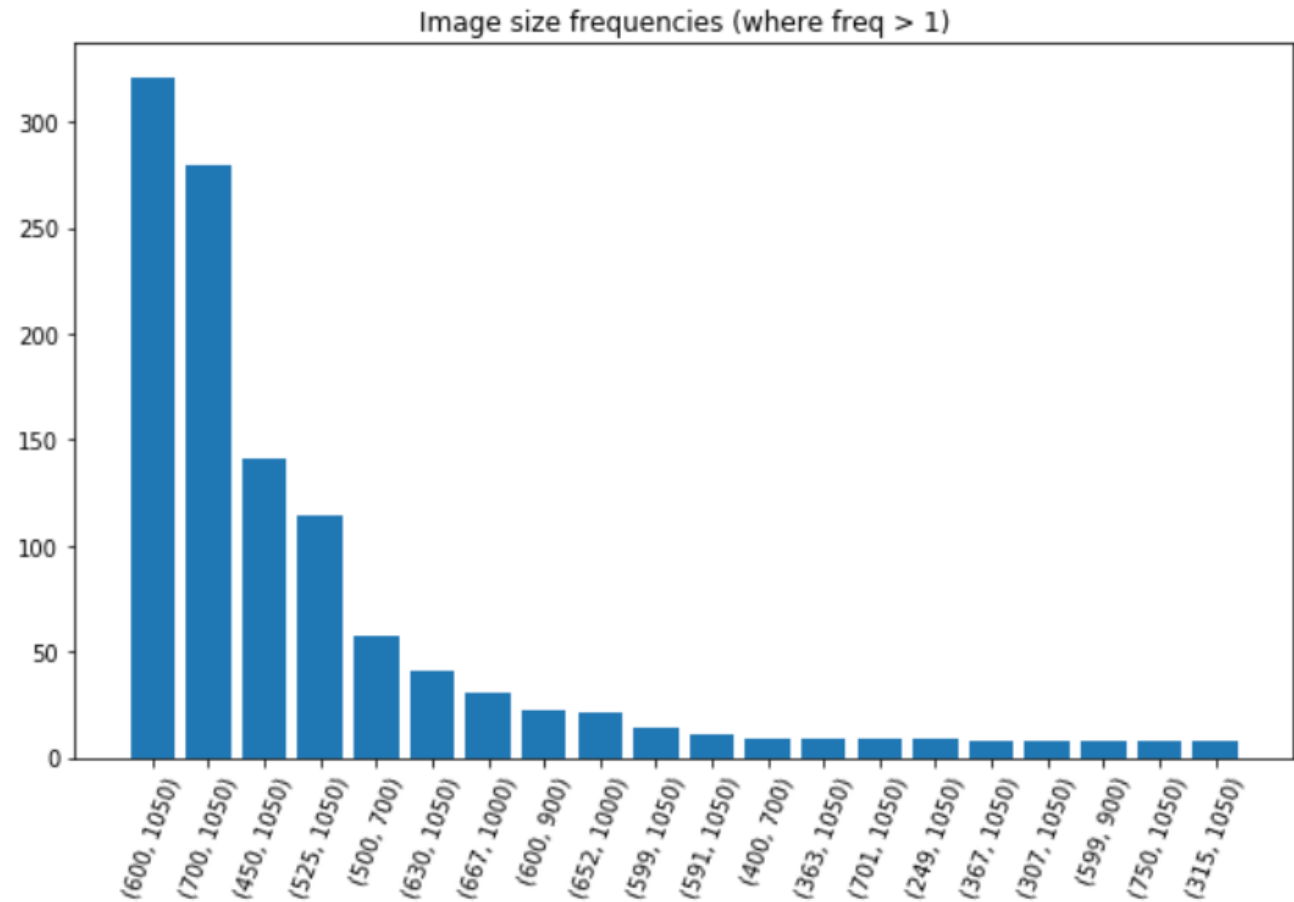
Problems

- ▶ The dataset includes a catch-all category called “new whale” which includes the vast majority of the images
- ▶ Technically speaking, you could just classify all images as new whale and be correct roughly 8% of the time
- ▶ A lot of classes contained only 1 image



More Problems

- ▶ Images varied greatly in size from 1050x1574 to 64x30
- ▶ Needed to somehow standardize size



Even More Problems

- ▶ Quite a few of the images were duplicated and placed either:
 - ▶ under the same label
 - ▶ Placed under several labels
 - ▶ placed under “new whale” in addition to other labels
- ▶ Found by translating images into hashes and grouping

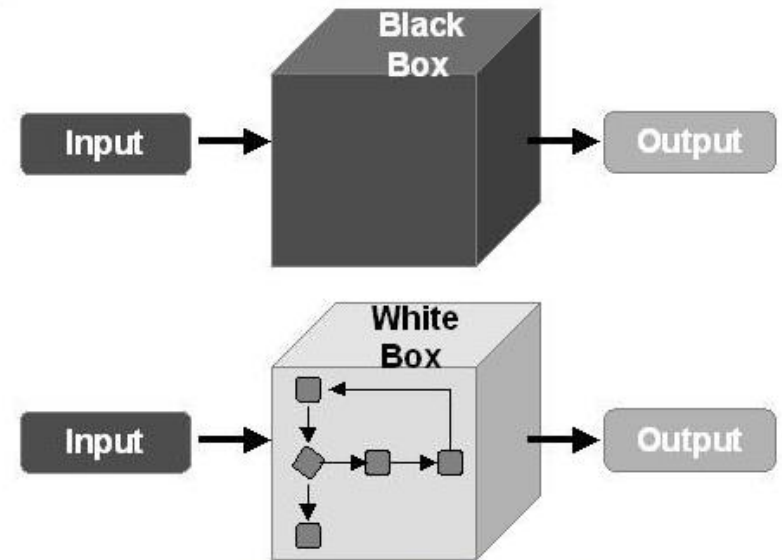
	Hash	Ids	Ids_count	Ids_contain_new_whale
bb8ec43039cb663c	3 {w_cae7677, new_whale}		2	True
9e1bc5d0bc4e0bc3	2 {new_whale, w_7185713}		2	True
af8fd0fcd3702940	2 {new_whale, w_1f09cdd}		2	True
e9889673ed9d5a02	2 {new_whale, w_a365757}		2	True
ee9ac1f47a4b8470	2 {new_whale, w_17a2610}		2	True
ed088dab92f0e3f0	2 {new_whale, w_ab4cae2}		2	True
932c5ac3a4b9ac5b	2 {w_2f54c3c}		1	False
8b90a4633b5cd29f	2 {w_dcb1f2a}		1	False
84717a9ec1a4717d	2 {w_ee948c6}		1	False
96bd2dc2c8c272f4	2 {w_5ba417d, new_whale}		2	True

Solutions for Data Preparation

- ▶ Removed the new whale category entirely
- ▶ Removed category's with < 10 images
- ▶ Removed duplicates from remainder
- ▶ Result was 48 classes which we padded with augmented data (over 2,000 images)

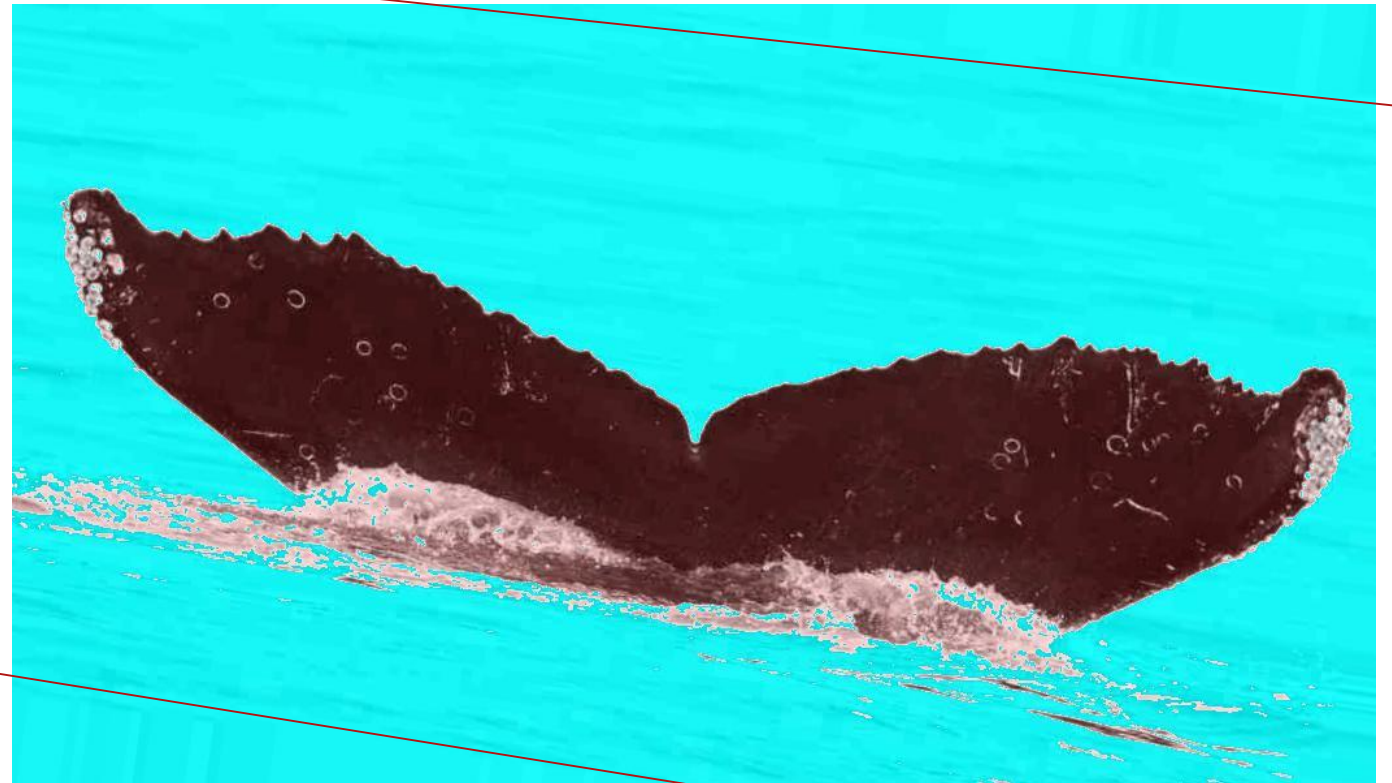
Classification Methods

- ▶ White Box Classification
 - ▶ SVM, KNN, Decision Tree, etc.
 - ▶ Needs feature extraction
 - ▶ Needs more data processing
 - ▶ More easy to be implemented
- ▶ Black Box
 - ▶ Deep Neural Network
 - ▶ No manual data extraction
 - ▶ Multiple feature methods applied automatically
 - ▶ Extremely resource intensive



Data Augmentation

- ▶ Set to high contrast and then applied greyscale
- ▶ Applied random rotations and shifts to create additional data
- ▶ Found a mean size then resized to 100 x 100 while maintaining aspect ratio
- ▶ Filled remaining area with random noise based on nearest cells



L2-Histogram of Gradients (HOG)

- ▶ HOG convolves the images with two filters that are sensitive to horizontal and vertical brightness gradients-allowing us to capture
 - ▶ Edge
 - ▶ Contour
 - ▶ Texture information
- ▶ HOG also subdivides the images into cells of a predetermined size, and computes histograms of the gradient orientations within each cell.
- ▶ Data is normalized in each cell and the result is a one dimensional feature vector made from the information in each cell of an image.
- ▶ Numpy array of (2256, 8100)



White Box Classifiers

- ▶ We started by implementing 5 basic classifiers without parameter tuning
 - ▶ K Nearest Neighbors
 - ▶ SVC (C-Support SVM). C is penalty value
 - ▶ NuSVC: similar to SVC but uses a parameter to control the number of support vectors
 - ▶ Decision Tree
 - ▶ Random Forest (messed up... supposed to use with decision tree)

```
=====
KNeighborsClassifier
****Results****
Accuracy: 19.9409%
Log Loss: 17.849415387393172
=====

SVC
****Results****
Accuracy: 41.8021%
Log Loss: 3.1330715145841355
=====

NuSVC
****Results****
Accuracy: 6.6470%
Log Loss: 3.7058826960689952
=====

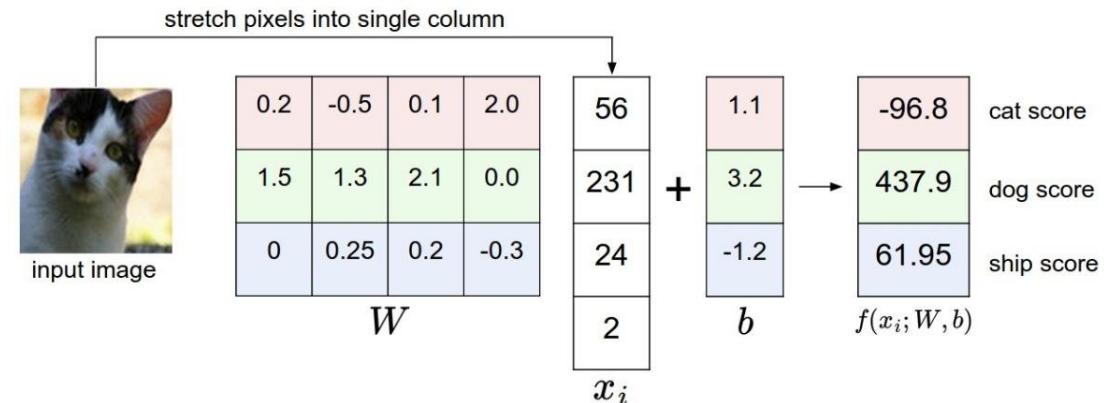
DecisionTreeClassifier
****Results****
Accuracy: 10.6352%
Log Loss: 30.865523957047262
=====

RandomForestClassifier
****Results****
Accuracy: 14.3279%
Log Loss: 19.576194929467743
=====
```

Parameter Tuning of SVM

► SVM:

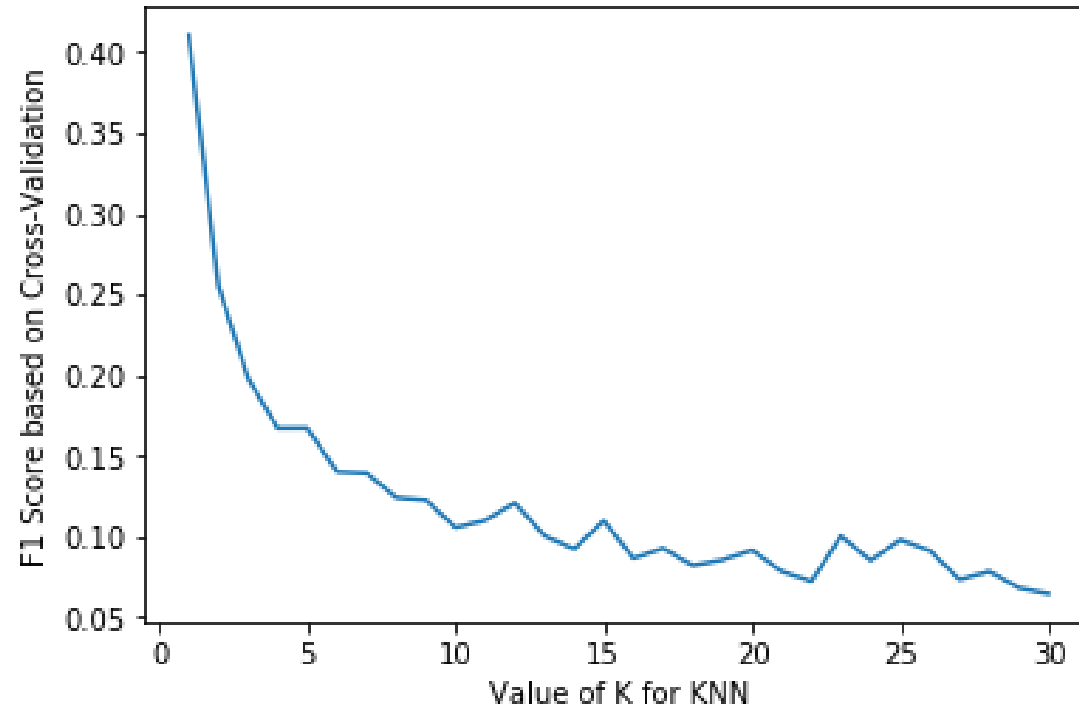
- Tried both rbf and linear through Grid Search
- Best Para: $C = 100$, linear kernel
- Linear kernel: the images are stretched into high-dimensional column (8100), each image is a single point in this space. Linear Classifiers computes the score of a class as a weighted sum of all of its pixels values across the column.



Result of SVM

- ▶ F1-Score avg 0.40
- ▶ Precision avg 0.44
- ▶ Recall avg 0.41
- ▶ For 48 classes using a linear kernel it's not horrible but not good

	precision	recall	f1-score
0	0.44	0.69	0.54
1	0.33	0.39	0.36
2	0.22	0.20	0.21
3	0.40	0.33	0.36
4	0.46	0.69	0.55
5	0.25	0.08	0.12
6	0.29	0.38	0.33
7	0.39	0.54	0.45
8	0.50	0.39	0.44
9	0.20	0.08	0.12
10	0.45	0.71	0.56
11	0.80	0.80	0.80
12	0.63	0.80	0.71
13	0.27	0.31	0.29
14	0.33	0.30	0.32
15	0.30	0.73	0.42
16	0.28	0.24	0.26
17	0.33	0.25	0.29
18	0.45	0.38	0.41
19	0.75	0.27	0.40
20	0.46	0.40	0.43
21	1.00	0.17	0.29
22	0.70	0.58	0.64
23	0.25	0.29	0.27
24	0.43	0.25	0.32
25	0.54	0.54	0.54
26	0.43	0.30	0.35
27	0.50	0.14	0.22
28	0.43	0.53	0.48
29	0.62	0.67	0.64
30	0.47	0.53	0.50
31	0.45	0.53	0.49
32	0.25	0.27	0.26
33	0.40	0.20	0.27
34	0.40	0.17	0.24
35	0.67	0.27	0.38
36	0.50	0.15	0.24
37	0.67	0.55	0.60
38	0.31	0.31	0.31
39	0.42	0.29	0.34
40	0.14	0.29	0.19
41	0.08	0.10	0.09
42	0.45	0.45	0.45
43	0.33	0.46	0.39
44	0.53	0.62	0.57
45	1.00	0.27	0.43
46	0.56	0.50	0.53
47	0.40	0.36	0.38
avg / total	0.44	0.41	0.40



Parameter Tuning of KNN

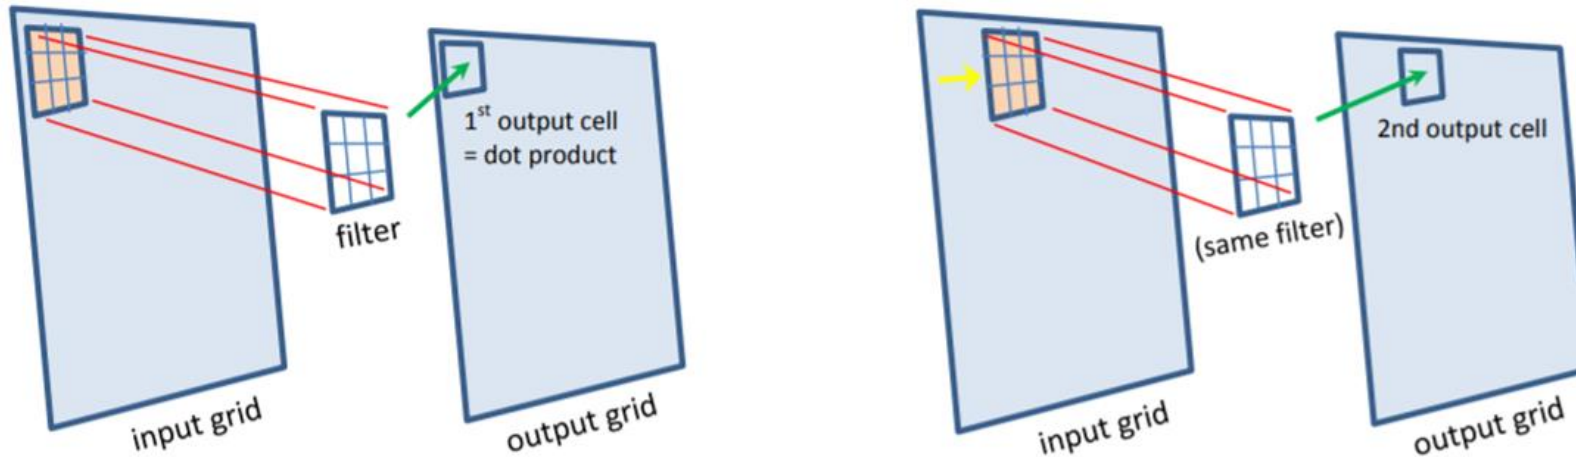
KNN Tuning Results

- ▶ F1-Score avg 0.43
- ▶ Precision avg 0.48
- ▶ Recall avg 0.43
- ▶ Better...

	precision	recall	f1-score
0	0.67	0.54	0.60
1	0.71	0.28	0.40
2	0.57	0.40	0.47
3	0.70	0.58	0.64
4	0.23	0.50	0.31
5	0.47	0.58	0.52
6	0.62	0.38	0.48
7	0.18	0.54	0.27
8	0.46	0.33	0.39
9	0.31	0.33	0.32
10	0.53	0.57	0.55
11	0.67	0.40	0.50
12	0.40	0.40	0.40
13	0.22	0.31	0.26
14	0.56	0.50	0.53
15	0.22	0.45	0.29
16	0.32	0.29	0.30
17	0.29	0.17	0.21
18	0.35	0.29	0.32
19	0.40	0.18	0.25
20	0.26	0.40	0.32
21	0.40	0.17	0.24
22	1.00	0.42	0.59
23	0.91	0.71	0.80
24	0.33	0.42	0.37
25	0.44	0.54	0.48
26	0.40	0.40	0.40
27	0.46	0.43	0.44
28	0.44	0.58	0.50
29	0.71	0.83	0.77
30	0.73	0.53	0.62
31	0.50	0.47	0.49
32	0.30	0.20	0.24
33	0.44	0.40	0.42
34	0.67	0.15	0.25
35	0.50	0.47	0.48
36	0.56	0.42	0.48
37	0.45	0.45	0.45
38	0.67	0.46	0.55
39	0.64	0.53	0.58
40	0.24	0.36	0.29
41	0.60	0.60	0.60
42	0.64	0.45	0.53
43	0.25	0.38	0.30
44	0.78	0.54	0.64
45	0.43	0.27	0.33
46	0.50	0.30	0.37
47	0.28	0.45	0.34
avg / total	0.48	0.43	0.43

Convolutional Layer

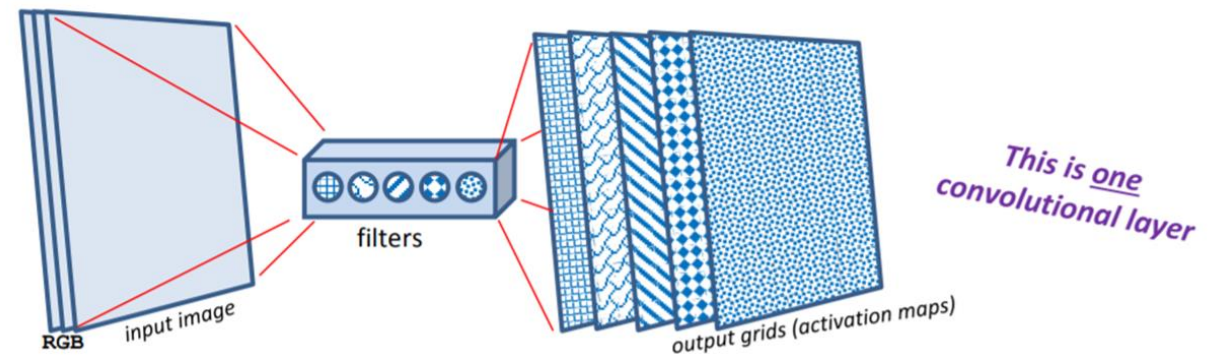
Typically used for input data that is organized into square 2D grids (such as for image recognition). A convolutional layer has a set of some number of “filters”. Each filter performs a series of dot products:

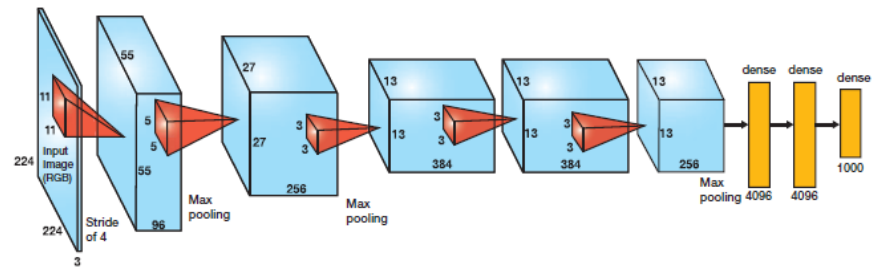
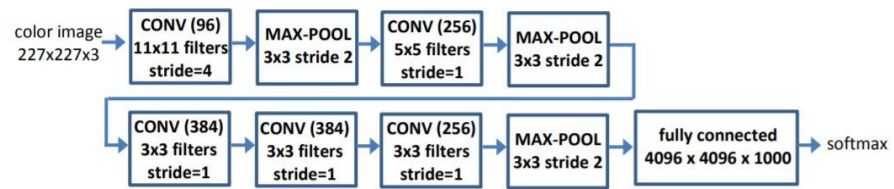


Convolutional Neural Network

CNN

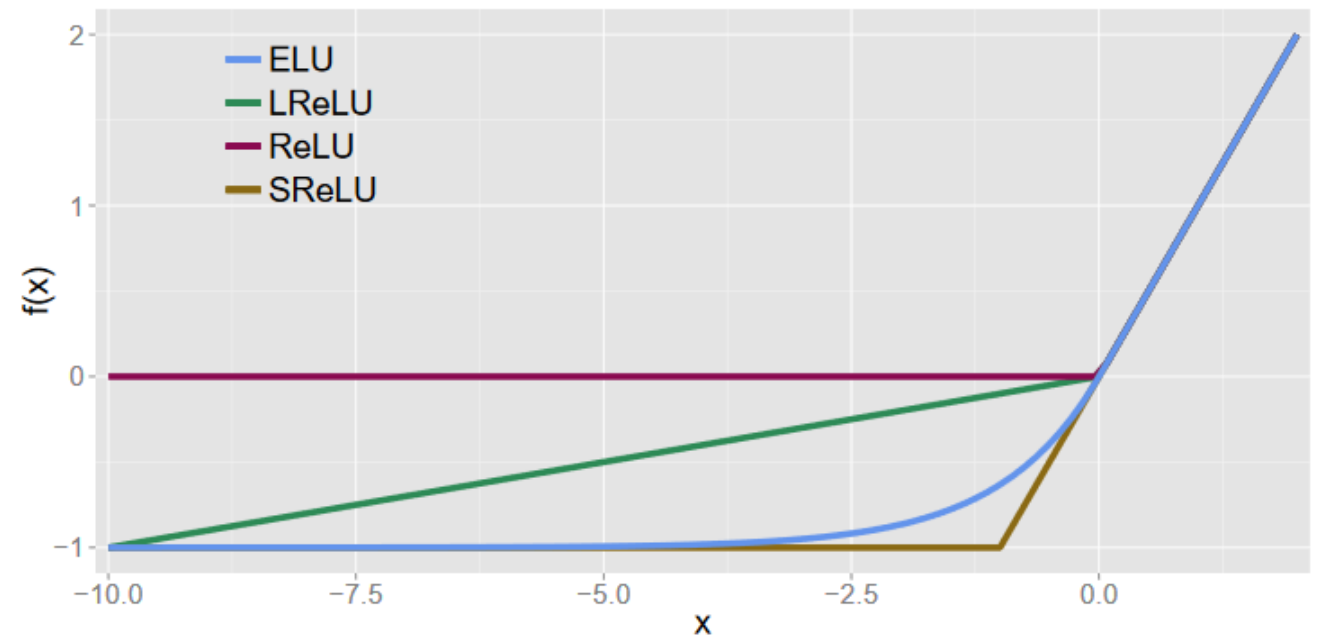
- ▶ Each filter extract one feature
- ▶ Early stage detect low level features
- ▶ Train on themselves





Modifications

- ▶ Based on paper:
<https://arxiv.org/pdf/1511.07289.pdf>
- ▶ Use Exponential Linear Units instead of Rectified ELU (ReLU)
- ▶ Add extra Layer
- ▶ Supposedly good for smaller datasets



```
1 import warnings
2 warnings.filterwarnings('ignore')
3
4 import tflearn
5 from tflearn.layers.core import input_data, dropout, fully_connected
6 from tflearn.layers.conv import conv_2d, max_pool_2d
7 from tflearn.layers.estimator import regression
8 from tflearn.metrics import Accuracy
9
10 acc = Accuracy()
11 network = input_data(shape=[None, 100, 100, 3])
12
13 # Conv layers -----
14 network = conv_2d(network, 64, 3, strides=1, activation='elu')
15 network = max_pool_2d(network, 2, strides=2)
16 network = conv_2d(network, 64, 3, strides=1, activation='elu')
17 network = max_pool_2d(network, 2, strides=2)
18 network = conv_2d(network, 64, 3, strides=1, activation='elu')
19 network = conv_2d(network, 64, 3, strides=1, activation='elu')
20 network = conv_2d(network, 64, 3, strides=1, activation='elu')
21 network = max_pool_2d(network, 2, strides=2)
22
23 # Fully Connected Layers -----
24 network = fully_connected(network, 1024, activation='tanh')
25 network = dropout(network, 0.5)
26 network = fully_connected(network, 1024, activation='tanh')
27 network = dropout(network, 0.5)
28
29 # CHANGE BELOW FOR NUMBER OF CLASSES
30 network = fully_connected(network, 57, activation='softmax')
31 network = regression(network, optimizer='momentum', loss='categorical_crossentropy',
32 learning_rate=0.001, metric=acc)
33
34 model = tflearn.DNN(network, tensorboard_verbose=3, tensorboard_dir="logs")
```

Results

- ▶ 2 hours to run on a Nvidia 1050ti
- ▶ F1-Score avg 0.89
- ▶ Precision avg 0.89
- ▶ Recall avg 0.89

	precision	recall	f1-score
0	0.87	0.79	0.83
1	0.83	0.91	0.87
2	0.88	0.95	0.92
3	1.00	0.82	0.90
4	0.91	0.83	0.87
5	0.72	1.00	0.84
6	1.00	0.95	0.97
7	0.90	0.64	0.75
8	0.89	1.00	0.94
9	0.91	0.71	0.80
10	1.00	0.71	0.83
11	0.93	1.00	0.96
12	0.85	0.89	0.87
13	1.00	0.67	0.80
14	0.94	0.84	0.89
15	0.87	0.87	0.87
16	1.00	0.73	0.84
17	0.79	0.88	0.84
18	0.88	0.91	0.89
19	0.92	0.92	0.92
20	0.86	0.89	0.87
21	1.00	0.85	0.92
22	0.84	0.94	0.89
23	1.00	0.69	0.82
24	1.00	0.54	0.70
25	1.00	0.94	0.97
26	1.00	0.92	0.96
27	0.81	0.81	0.81
28	0.92	0.92	0.92
29	0.88	0.88	0.88
30	0.91	0.91	0.91
31	1.00	0.92	0.96
32	0.85	0.92	0.88
33	0.92	0.73	0.81
34	0.93	0.78	0.85
35	0.85	0.81	0.83
36	0.93	0.76	0.84
37	1.00	0.79	0.88
38	1.00	0.91	0.95
39	0.79	1.00	0.88
40	0.83	0.88	0.86
41	0.75	0.80	0.77
42	0.87	0.93	0.90
43	0.90	0.64	0.75
44	0.88	0.71	0.79
45	1.00	0.91	0.95
46	0.93	0.72	0.81
47	0.91	0.91	0.91
48	0.92	1.00	0.96
49	0.86	0.55	0.67
50	0.88	0.88	0.88
51	0.88	0.94	0.91
52	0.85	0.85	0.85
53	1.00	0.67	0.80
54	0.91	0.91	0.91
55	1.00	0.69	0.82
56	0.90	0.75	0.82
avg / total	0.89	0.89	0.89

Conclusions

- ▶ CNN Provide best results for image classification on this scale
- ▶ SVM and KNN apply a SoftMax and are essentially just one layer of CNN
- ▶ CNN takes a very long time even with GPU
- ▶ CNN doesn't reveal features extracted, however more are extracted
- ▶ BE CONSISTENT WITH LIBRARIES
- ▶ If doing images, don't do what we did. Just do CNN and maybe apply SVM at the end to improve performance